

# Fast Serialization of Numpy Arrays with Bloscpack

Valentin Haenel

Freelance Consultant and Software Developer

@esc\_

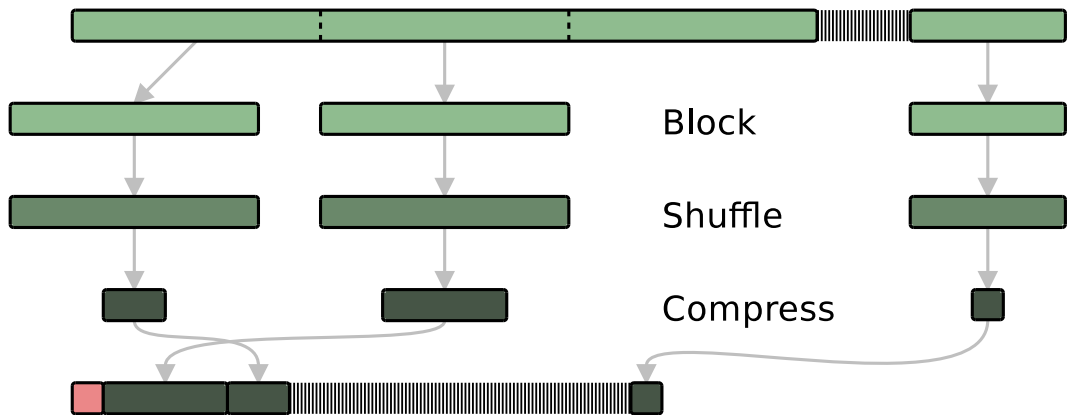
27 July 2014 - PyData Berlin

Version: 2014-PyDataBerlin    <https://github.com/esc/PyDataBerlin2014-bloscpack>  
This work is licensed under the *Creative Commons Attribution-ShareAlike 3.0 License*.

Blosc


## Blosc: A Fast Meta-Codec

- ▶ Blocking
- ▶ Shuffling
- ▶ Multithreaded
- ▶ Multi-codec



## Shuffle Filter

- ▶ Reorder bytes by significance inside a block
- ▶ Potentially reduce Lempel-Ziv complexity of the data



# Multi-Codec

- ▶ By default it uses **Blosclz** – derived from **Fastlz**
- ▶ Alternative codecs
  - ▶ **LZ4 / LZ4HC**
  - ▶ **Snappy**
  - ▶ **Zlib**

## python-blosc: Bindings

- ▶ Python C-API bindings
- ▶ Accepts a pointer as int

Bloscpack


# Bloscpack

- ▶ Simple serialization format based on Blosc
- ▶ Command line interface
- ▶ Python API with support for Numpy arrays
- ▶ Aimed at developers / power-users



# Features

- ▶ Chunked, compressed format
- ▶ Metadata (optional)
- ▶ Checksums (optional)
- ▶ Offsets, including pre-allocation for append (optional)



# Use Cases

- ▶ Fast serialization
- ▶ Streaming
- ▶ On disk columnar storage
- ▶ Substrate on which to build high-level abstractions

# Python API

- ▶ Inner loop compress/decompress implemented
- ▶ Concept: sinks and sources
  - ▶ PlainSource -> CompressedSink
  - ▶ CompressedSource -> PlainSink
- ▶ Supply appropriate source and sink
- ▶ Sources and sinks must obey an interface/contract
- ▶ Get easy Anything -> Anything
- ▶ E.g. Numpy -> {string, file, memory, network}

## pack and unpack

```
def pack(source, sink,
        nchunks, chunk_size, last_chunk,
        metadata=None,
        blosc_args=None,
        bloscpack_args=None,
        metadata_args=None):
    pass

def unpack(source, sink):
    pass
```

## Numpy Example

```
import numpy as np
import bloscpack as bp

a = np.arange(1e7)

# pack with defaults
bp.pack_ndarray_file(a, 'a.blp')
```

## Numpy Example

```
# pack with custom settings  
bp.pack_ndarray_file(a, 'a.blp',  
    chunk_size='20M',  
    blosc_args=bp.BloscArgs(cname='lz4', clevel=9),  
    bloscpack_args=bp.BloscpackArgs(offsets=False),  
    )
```

## Numpy Example

```
# unpack
```

```
b = bp.unpack_ndarray('a.blp')
```

## Commandline Example

```
$ blpk compress --level 9 --codec lz4 --no-offsets data.dat
```



## Extension Example

- ▶ Idea: how about S3 connectivity?
- ▶ Implement CompressedS3Sink and CompressedS3Source
- ▶ (These know nothing about Numpy)
- ▶ Result: ability to compress a Numpy array to an S3 bucket

## Relationship to (Distributed) Analytics Engines

- ▶ Column-oriented, compressed, chunked storage
  - ▶ bcolz
  - ▶ Hustle
  - ▶ Parquet
  - ▶ RCFile / ORCFile
- ▶ Fast, partial loading from disk or network
- ▶ Reduced storage requirements
- ▶ But: need to chose the *right codec*<sup>TM</sup>
- ▶ A Bloscpack file translates directly to a serialized column

## Benchmarks

## Background

- ▶ Builds on benchmarks presented at EuroScipy 2013
- ▶ Those used a laptop with SSD and SD storage
- ▶ Showed that Bloscpack can be outperform contenders

See also: Bloscpack: a compressed lightweight serialization format for numerical data

# Experimental Setup

- ▶ Use Python 3.4
- ▶ Use some real-world datasets
- ▶ Benchmark new codecs available in Blosc
- ▶ Add PyTables to the mix
- ▶ Run it in the AWS cloud

# Datasets

- ▶ **arange**
  - ▶ Integers
- ▶ **linspace**
  - ▶ floats
- ▶ **poisson**
  - ▶ more or less random numbers
- ▶ **neuronal**
  - ▶ Neural net spike time stamps
  - ▶ Kindly provided by Yuri Zaytsev
- ▶ **bitcoin**
  - ▶ Historical MtGOX trade data

# Contenders

- ▶ PyTables
  - ▶ HDF5 interface
  - ▶ Supports Blosc and others
- ▶ NPY
  - ▶ Numpy plain serialization
- ▶ NPZ
  - ▶ Numpy compressed (using zip) serialization
- ▶ ZFile
  - ▶ Joblib's compressed (using zlib) **pickler** extension

## NPY Flaw

- ▶ Prior to serialization, array is copied in memory with `tostring()`
- ▶ Fixed by Olivier Grisel to use `nditer` (#4077)
- ▶ Available in v1.9.0b1, which is what I used for the benchmarks



## NPZ Flaw

- ▶ Create a temporary plain version (/tmp)
- ▶ Compresses into a Zip archive from there
- ▶ Due to issues with the ZipFile module

## ZFile Flaw

- ▶ Does not support arrays larger than 2GB
- ▶ An `int32` is used somewhere for the size in the `zlib` module

# Remaining Experimental Parameters

- ▶ Instance
  - ▶ c3.2xlarge
  - ▶ CPUs: 8
  - ▶ RAM: 15GB
- ▶ Dataset Sizes
  - ▶ 1MB
  - ▶ 10MB
  - ▶ 100MB
- ▶ Storage
  - ▶ EBS
  - ▶ Ephemeral

# Measurements

- ▶ Writing to disk is tricky
- ▶ Measure with hot and cold FS cache
- ▶ Add disk sync to the timing
- ▶ Used a variant to the `timeit` utility.

## Results

Let's look at the `arange` and `neuronal` datasets in the `small` and `large` configuration on `ebs` → IPython notebooks

## Aggregated Results

- ▶ Single plots can supply insights
- ▶ Need to aggregate for a big picture
- ▶ Award points to a codec/level combination
  - ▶ Slowest receives 1 point
  - ▶ Fastest receives max points
  - ▶ Compute with and without ratio
- ▶ Recommendation for a good general purpose codec
- ▶ See → Ipython notebook aggregate

## Conclusions – What did I Observe?

- ▶ Bloscpack vs. plain
  - ▶ In general it will not hurt to try Bloscpack
- ▶ Bloscpack vs. NPZ/ZFile
  - ▶ These formats don't scale well to large arrays
- ▶ Bloscpack vs. PyTables
  - ▶ Bloscpack is somewhat better at fast serialization
  - ▶ PyTables isn't the worst choice for long-term storage – but do use Blosc
- ▶ Blosclz vs. LZ4 vs. LZ4HC vs. Snappy vs. Zlib
  - ▶ Blosclz and LZ4 are the kings of fast compression
  - ▶ Snappy seems pretty average
  - ▶ Zlib can really benefit from Blosc acceleration and shuffle

# Reproducibility

- ▶ Results contained in the talk sources repository
- ▶ Lists almost all the hashes and configurations
- ▶ All code open source
- ▶ All datasets additionally available from backup location on own infrastructure
- ▶ AMI available incl. instructions (soon to come / ask me)



# TODO

- ▶ Find other ways to analyse the results
- ▶ Stabilize the format
- ▶ Release Python 3 support
- ▶ Support Bloscpack in Joblib
  - ▶ Speed gain
  - ▶ Mitigate 2GB issue

## Getting In Touch

`http://blosc.org`

`@esc_`